

Central test LRegression

March 20, 2024

Ordinary linear regression

```
[1]: #importing libraries
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
```

```
[2]: data = pd.read_csv('C:\\Users\\DELL PC\\Desktop\\LINEAR PROG\\Ice Cream Sales -_
↳temperatures.csv')
data
```

```
[2]:
```

	Temperature	Ice Cream Profits
0	39	13.17
1	40	11.88
2	41	18.82
3	42	18.65
4	43	17.02
..
360	99	85.13
361	99	87.08
362	99	89.29
363	101	81.91
364	101	85.02

[365 rows x 2 columns]

```
[3]: temperature_list=data['Temperature'].tolist()
ice_cream_profits=data['Ice Cream Profits'].tolist()
x = np.array(temperature_list, dtype = np.float64).reshape(-1,1)
y = np.array(ice_cream_profits, dtype = np.float64)
x
```

```
[3]: array([[ 39.],
           [ 40.],
           [ 41.],
           [ 42.],
           [ 43.],
           [ 43.]])
```

[44.],
[44.],
[45.],
[45.],
[45.],
[46.],
[46.],
[47.],
[48.],
[48.],
[48.],
[48.],
[48.],
[48.],
[49.],
[49.],
[50.],
[50.],
[50.],
[50.],
[50.],
[51.],
[51.],
[52.],
[52.],
[52.],
[52.],
[52.],
[53.],
[53.],
[53.],
[53.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[55.],
[55.],
[55.],
[55.],
[56.],
[56.],

[57.],
[57.],
[58.],
[58.],
[58.],
[58.],
[58.],
[59.],
[59.],
[59.],
[59.],
[59.],
[59.],
[59.],
[60.],
[60.],
[60.],
[60.],
[60.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[62.],
[62.],
[62.],
[62.],
[63.],
[63.],
[63.],
[63.],
[63.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],

[64.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[67.],
[67.],
[67.],
[67.],
[67.],
[67.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[69.],
[69.],
[69.],
[69.],
[69.]

[69.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],

[illegible]

[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[79.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[80.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[81.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[82.],
[83.],
[83.],
[83.],
[83.],
[83.],
[83.],

[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[84.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[87.],
[87.],
[87.],
[87.],
[87.],
[87.],
[88.],
[88.],
[88.],
[89.],
[89.],
[89.],
[89.],
[89.],
[90.],
[90.],


```

[ 90.],
[ 90.],
[ 90.],
[ 90.],
[ 90.],
[ 91.],
[ 91.],
[ 91.],
[ 92.],
[ 92.],
[ 92.],
[ 92.],
[ 93.],
[ 93.],
[ 94.],
[ 94.],
[ 95.],
[ 95.],
[ 95.],
[ 95.],
[ 96.],
[ 96.],
[ 96.],
[ 97.],
[ 98.],
[ 99.],
[ 99.],
[ 99.],
[101.],
[101.]]

```

```
[4]: y
```

```

[4]: array([13.17, 11.88, 18.82, 18.65, 17.02, 15.88, 19.07, 19.57, 21.62,
          22.34, 19.23, 21.25, 19.81, 22.12, 24.22, 24.68, 23.78, 26.41,
          25.01, 22.29, 27.81, 23.54, 22.89, 25.68, 27.29, 27.64, 27.31,
          21.93, 32.18, 30.67, 28.05, 28.82, 27.87, 29.39, 32.6 , 31.62,
          25.71, 28.48, 30.09, 33.58, 29.75, 31.94, 33.71, 28.37, 27.41,
          27.99, 30.37, 27.68, 29.53, 33.91, 34.19, 33.22, 34.47, 30.89,
          35.8 , 33.44, 36.79, 31.56, 35.13, 36.11, 32.39, 38.18, 29.69,
          38.47, 37.74, 36.71, 32.29, 37.5 , 35.33, 35.06, 36.25, 40.25,
          39.69, 40.95, 37.96, 38.1 , 38.21, 37.3 , 39.53, 37.42, 39.42,
          38.16, 37.66, 39.04, 41.44, 40.19, 37.93, 50.17, 44.15, 41.58,
          40.59, 39.17, 40.57, 40.28, 41.21, 44.85, 40.94, 40.14, 38.57,
          44.07, 44.1 , 47.36, 45.38, 41.09, 43.78, 42.72, 42.1 , 43.28,
          44.31, 42.71, 43.03, 42.16, 46.74, 47.68, 44.48, 47.52, 44.98,
          45.07, 45.42, 47.36, 48.26, 51.75, 45.05, 40.65, 48.65, 45.26,

```

```

46.04, 44.85, 42.94, 50.62, 45.65, 49.37, 45.89, 50.74, 47.17,
49.6 , 41.68, 46.9 , 47.35, 47.73, 43.73, 47.47, 51.38, 41.74,
49.88, 47.78, 42.5 , 48.77, 49.46, 50.87, 49.12, 49.95, 50.31,
49.32, 52.67, 52.05, 48.82, 53.33, 54.59, 53.77, 49.6 , 52.17,
46.74, 53.04, 49.34, 55.04, 57.18, 51.26, 53.78, 51.55, 50.01,
53.59, 52.47, 48.96, 53.57, 50.79, 52.13, 52.42, 54.67, 51.82,
53.21, 54.4 , 55.01, 54.08, 53.97, 55.28, 54.36, 53.62, 50.65,
55.52, 58.61, 50.64, 54.28, 53.95, 53.44, 57.1 , 54.26, 55.34,
53.71, 57.84, 55.91, 58.62, 58.85, 52.84, 56.59, 59.43, 59.69,
53.83, 59.41, 53.17, 53.48, 59.94, 60.31, 60.33, 53.82, 53.07,
59.48, 54.1 , 56.33, 59.87, 60.75, 56.43, 60.86, 55.07, 58.39,
58.72, 57.52, 56.33, 57.47, 58.13, 60.46, 60.33, 60.89, 62.58,
61.22, 59.62, 58.31, 59.12, 57.93, 57.25, 62.2 , 59.7 , 64.82,
57.06, 62.52, 59.93, 61.71, 59.49, 67.42, 56.34, 59.69, 57.44,
64.63, 55.47, 61.22, 62.79, 59.91, 61.59, 63.46, 64.45, 65.42,
61.82, 64.36, 58.11, 59.47, 65.86, 61.52, 62.12, 64.23, 62.36,
62.32, 64.97, 66.15, 64.02, 63.41, 61.85, 65.49, 64.39, 66.06,
64.86, 62.85, 66.57, 65.54, 62.58, 63.29, 64.38, 60.78, 65.66,
66.61, 65.12, 63.13, 63.35, 65.4 , 65.41, 68.28, 64.1 , 66.26,
63.63, 67.58, 68.54, 65.2 , 67.93, 67.88, 69.71, 64.22, 61.82,
68.28, 62.99, 64.96, 65.99, 70.3 , 64.31, 69.59, 68.35, 69.66,
71.46, 69.9 , 69.19, 67.97, 64.85, 70.43, 68.48, 70.29, 65.19,
68. , 70.64, 69.67, 74.69, 69.78, 73.16, 71.51, 73.32, 74.09,
71.12, 67.58, 77.39, 75.11, 74.8 , 73.94, 75.94, 79.31, 81.81,
75.58, 78.2 , 75.6 , 75.04, 77.41, 79.76, 77.18, 80.94, 75.7 ,
78.2 , 80.75, 80.97, 80.98, 80.02, 82.83, 80.95, 82.5 , 84.12,
85.13, 87.08, 89.29, 81.91, 85.02])

```

```

[5]: #checking the missing data
data.isna().sum()

```

```

[5]: Temperature      0
Ice Cream Profits    0
dtype: int64

```

```

[6]: #Splitting the dataset
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, train_size=0.8)

```

```

[7]: #create the model and the fit the data in it
model = LinearRegression()
model.fit(x_train, y_train)

```

```

[7]: LinearRegression()

```

```

[8]: #make predictions
y_pred = model.predict(x_test)

```

```
y_pred
```

```
[8]: array([59.42452428, 47.48502099, 76.13982889, 35.54551769, 57.03662362,
         54.64872296, 28.38181571, 58.23057395, 57.03662362, 39.12736868,
         43.90317    , 59.42452428, 63.00637527, 49.87292164, 16.44231242,
         43.90317    , 70.17007725, 65.39427593, 76.13982889, 76.13982889,
         23.6060144 , 79.72167988, 67.78217659, 67.78217659, 59.42452428,
         55.84267329, 72.5579779 , 24.79996472, 34.35156736, 57.03662362,
         58.23057395, 67.78217659, 14.05441176, 39.12736868, 60.61847461,
         60.61847461, 27.18786538, 65.39427593, 61.81242494, 65.39427593,
         67.78217659, 25.99391505, 76.13982889, 59.42452428, 47.48502099,
         64.2003256 , 58.23057395, 51.06687197, 17.63626275, 53.45477263,
         45.09712033, 68.97612692, 64.2003256 , 64.2003256 , 45.09712033,
         39.12736868, 39.12736868, 52.2608223 , 61.81242494, 48.67897131,
         59.42452428, 29.57576604, 41.51526934, 58.23057395, 45.09712033,
         52.2608223 , 66.58822626, 36.73946802, 46.29107066, 67.78217659,
         58.23057395, 33.15761703, 66.58822626])
```

```
[9]: #Model accuracy
General_score = model.score(x_test,y_test)
model_score=model.score(x_train,y_train)
```

```
[10]: #Evaluating the model
#model.coef_      -to find the coefficient/ model.intercept_ -to find the_
↪intercept
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
MAE = mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print('r2                : ',r2)
print('mean_absolute_error:',MAE)
print('mean_squared_error : ',MSE)
print('General_score      : ',General_score)
print('model_score        : ',model_score)
```

```
r2                : 0.9763452606634333
mean_absolute_error: 1.8914456150513512
mean_squared_error : 5.840620626586573
General_score      : 0.9763452606634333
model_score        : 0.9770972572223515
```

OPTIMISED LINEAR REGRESSION

```
[11]: #import necessary libraries
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
```

```

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score

#Load the data
data = pd.read_csv('C:\\Users\\DELL PC\\Desktop\\LINEAR PROG\\Ice Cream Sales -_
↳temperatures.csv')
data

```

```

[11]:      Temperature  Ice Cream Profits
0           39           13.17
1           40           11.88
2           41           18.82
3           42           18.65
4           43           17.02
..          ...           ...
360          99           85.13
361          99           87.08
362          99           89.29
363         101           81.91
364         101           85.02

```

[365 rows x 2 columns]

```

[12]: #Split the dataset into training and testing
temperature_list=data['Temperature'].tolist()
ice_cream_profits=data['Ice Cream Profits'].tolist()
x = np.array(temperature_list, dtype = np.float64).reshape(-1,1)
y = np.array(ice_cream_profits, dtype = np.float64)
x

```

```

[12]: array([[ 39.],
              [ 40.],
              [ 41.],
              [ 42.],
              [ 43.],
              [ 43.],
              [ 44.],
              [ 44.],
              [ 45.],
              [ 45.],
              [ 45.],
              [ 46.],
              [ 46.],
              [ 47.],
              [ 48.],
              [ 48.],
              [ 48.]])

```

[48.],
[48.],
[48.],
[49.],
[49.],
[50.],
[50.],
[50.],
[50.],
[50.],
[51.],
[51.],
[52.],
[52.],
[52.],
[52.],
[52.],
[53.],
[53.],
[53.],
[53.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[54.],
[55.],
[55.],
[55.],
[55.],
[56.],
[56.],
[57.],
[57.],
[58.],
[58.],
[58.],
[58.],
[58.],
[58.],
[59.],
[59.],
[59.],

[59.],
[59.],
[59.],
[59.],
[59.],
[60.],
[60.],
[60.],
[60.],
[60.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[61.],
[62.],
[62.],
[62.],
[62.],
[63.],
[63.],
[63.],
[63.],
[63.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[64.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[65.],
[66.],

[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[66.],
[67.],
[67.],
[67.],
[67.],
[67.],
[67.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[68.],
[69.],
[69.],
[69.],
[69.],
[69.],
[69.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[70.],
[71.],

[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[71.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[72.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[73.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[74.],
[75.],
[75.],
[75.],
[75.],
[75.],
[75.],
[76.],
[76.],

[illegible]

[illegible]

[84.],
[84.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[85.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[86.],
[87.],
[87.],
[87.],
[87.],
[87.],
[87.],
[88.],
[88.],
[88.],
[89.],
[89.],
[89.],
[89.],
[90.],
[90.],
[90.],
[90.],
[90.],
[90.],
[90.],
[90.],
[91.],
[91.],
[91.],
[92.],
[92.],
[92.],

```
[ 92.],
[ 93.],
[ 93.],
[ 94.],
[ 94.],
[ 95.],
[ 95.],
[ 95.],
[ 95.],
[ 96.],
[ 96.],
[ 96.],
[ 97.],
[ 98.],
[ 99.],
[ 99.],
[ 99.],
[101.],
[101.]])
```

```
[13]: x_train, x_test,y_train,y_test =train_test_split(x,y, train_size=0.8,
↳random_state=45)
```

```
[14]: #Performing the grid search & create a linear model
model = LinearRegression()
model
```

```
[14]: LinearRegression()
```

```
[15]: #Defining parameter grid search
param_grid = {
    'fit_intercept':[True, False],
    'copy_X':[True, False],
    'n_jobs':[None, True, False],
    'positive':[True, False],
}
param_grid
```

```
[15]: {'fit_intercept': [True, False],
      'copy_X': [True, False],
      'n_jobs': [None, True, False],
      'positive': [True, False]}
```

```
[16]: #performing the grid search
grid_search = GridSearchCV(model, param_grid, cv=5)
grid_search
```

```
[16]: GridSearchCV(cv=5, estimator=LinearRegression(),
                param_grid={'copy_X': [True, False],
                            'fit_intercept': [True, False],
                            'n_jobs': [None, True, False],
                            'positive': [True, False]})
```

```
[17]: #fitting the grid search into data
      grid_search.fit(x_train, y_train)
      grid_search
```

```
[17]: GridSearchCV(cv=5, estimator=LinearRegression(),
                param_grid={'copy_X': [True, False],
                            'fit_intercept': [True, False],
                            'n_jobs': [None, True, False],
                            'positive': [True, False]})
```

```
[18]: #finding the best_parameters
      best_params = grid_search.best_params_
      best_params
```

```
[18]: {'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'positive': True}
```

```
[19]: #Train the linear Regression model using best params
      best_model = LinearRegression(**best_params)
      best_model.fit(x_train, y_train)
```

```
[19]: LinearRegression(positive=True)
```

```
[20]: #Making predictions
      y_pred = best_model.predict(x_test)
      y_pred
```

```
[20]: array([50.91368179, 65.35438372, 52.11707362, 49.71028997, 54.52385728,
        17.21871064, 65.35438372, 37.6763717 , 37.6763717 , 58.13403276,
        43.69333083, 42.48993901, 82.2018693 , 24.4390616 , 68.9645592 ,
        25.64245343, 61.74420824, 19.62549429, 44.89672266, 76.18491016,
        52.11707362, 59.33742458, 67.76116737, 59.33742458, 66.55777555,
        32.86280439, 30.45602074, 44.89672266, 41.28654718, 67.76116737,
        58.13403276, 66.55777555, 36.47297987, 44.89672266, 55.7272491 ,
        47.30350631, 78.59169382, 68.9645592 , 47.30350631, 64.15099189,
        47.30350631, 16.01531881, 73.77812651, 50.91368179, 28.04923708,
        67.76116737, 31.65941256, 59.33742458, 84.60865295, 30.45602074,
        42.48993901, 38.87976352, 66.55777555, 52.11707362, 37.6763717 ,
        54.52385728, 53.32046545, 70.16795103, 53.32046545, 66.55777555,
        36.47297987, 54.52385728, 61.74420824, 58.13403276, 80.99847747,
        58.13403276, 46.10011449, 58.13403276, 66.55777555, 49.71028997,
        43.69333083, 64.15099189, 67.76116737])
```

```
[21]: #determining the accuracy of the model
MAE = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)

print('mean_absolute_error:',MAE)
print('r2 Score          : ',r2)
print('mean_squared_error : ',MSE)
```

```
mean_absolute_error: 1.7696053861131733
r2 Score           : 0.9787794320233247
mean_squared_error : 4.824166540507413
```

ORDINARY LOGISTIC MODEL

```
[22]: #importing libraries
import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression

data = pd.read_csv('C:\\Users\\DELL PC\\Desktop\\LINEAR PROG\\Logistic_dataset.
↳csv')
data
```

```
[22]:
```

	Feature 1	Feature 2	Label
0	1.764052	4.764052	0
1	0.978738	-2.021262	1
2	1.867558	4.867558	0
3	0.950088	-2.049912	1
4	-0.103219	2.896781	0
..
95	-1.292857	-4.292857	1
96	-0.039283	-3.039283	1
97	0.523277	-2.476723	1
98	0.771791	3.771791	0
99	2.163236	5.163236	0

[100 rows x 3 columns]

```
[23]: n = data.drop(['Label'], axis=1) #Cutting off the named column, then assigning
↳the rest of the data to n
n
```

```
[23]:
```

	Feature 1	Feature 2
0	1.764052	4.764052
1	0.978738	-2.021262
2	1.867558	4.867558
3	0.950088	-2.049912

```

4    -0.103219    2.896781
..          ...          ...
95   -1.292857   -4.292857
96   -0.039283   -3.039283
97    0.523277   -2.476723
98    0.771791    3.771791
99    2.163236    5.163236

```

[100 rows x 2 columns]

```
[24]: p = data['Label']
      p
```

```

[24]: 0    0
      1    1
      2    0
      3    1
      4    0
      ..
     95    1
     96    1
     97    1
     98    0
     99    0
      Name: Label, Length: 100, dtype: int64

```

```

[25]: #splitting data into training and testing dataset
      n_train, n_test, p_train, p_test = train_test_split(n,p, train_size=0.8,
      ↪random_state=45)

      #creating a model
      model = LogisticRegression()
      model.fit(n_train, p_train)

```

```
[25]: LogisticRegression()
```

```

[26]: #making predictions
      p_pred = model.predict(n_test)
      p_pred

```

```

[26]: array([0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0],
      dtype=int64)

```

```

[27]: #Evaluating the model
      from sklearn.metrics import accuracy_score,precision_score,f1_score,recall_score
      Accuracy_score = accuracy_score(p_test, p_pred)
      precision_score = precision_score(p_test, p_pred)

```

```
f1_score = f1_score(p_test,p_pred)
recall_score = recall_score(p_test, p_pred)
results = model.score(n_test, p_pred)

print('accuracy_score  : ',Accuracy_score)
print('precision_score  : ',precision_score)
print('f1_score         : ',f1_score)
print('recall_score     : ',recall_score)
print('results          : ',results)
```

```
accuracy_score  : 1.0
precision_score : 1.0
f1_score        : 1.0
recall_score    : 1.0
results         : 1.0
```

OPTIMISED LOGISTIC REGRESSION

```
[28]: #importing libraries
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, precision_score

#load the data
data = pd.read_csv('C:\\Users\\DELL PC\\Desktop\\LINEAR PROG\\Logistic_dataset.
↪csv')
data
```

```
[28]:
```

	Feature 1	Feature 2	Label
0	1.764052	4.764052	0
1	0.978738	-2.021262	1
2	1.867558	4.867558	0
3	0.950088	-2.049912	1
4	-0.103219	2.896781	0
..
95	-1.292857	-4.292857	1
96	-0.039283	-3.039283	1
97	0.523277	-2.476723	1
98	0.771791	3.771791	0
99	2.163236	5.163236	0

```
[100 rows x 3 columns]
```

```
[29]: n = data.drop(['Label'], axis=1) #Cutting off the named column, then assigning
↪the rest of the data to n
n
```



```
[29]:
```

	Feature 1	Feature 2
0	1.764052	4.764052
1	0.978738	-2.021262
2	1.867558	4.867558
3	0.950088	-2.049912
4	-0.103219	2.896781
..
95	-1.292857	-4.292857
96	-0.039283	-3.039283
97	0.523277	-2.476723
98	0.771791	3.771791
99	2.163236	5.163236

[100 rows x 2 columns]

```
[30]: p = data['Label']
p
```

```
[30]:
```

0	0
1	1
2	0
3	1
4	0
..	..
95	1
96	1
97	1
98	0
99	0

Name: Label, Length: 100, dtype: int64

```
[31]: #splitting data into training and testing dataset
n_train, n_test, p_train, p_test = train_test_split(n,p, train_size=0.8,
↳random_state=45)
```

```
[32]: #building the model
model = LogisticRegression()
model
```

```
[32]: LogisticRegression()
```

```
[35]: #fitting the model
model.fit(n_train,p_train)
```

```
[35]: LogisticRegression()
```

```
[36]: #making predictions
p_pred = model.predict(n_test)
p_pred
```

```
[36]: array([0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0],
      dtype=int64)
```

```
[43]: import warnings
warnings.filterwarnings("ignore")
#Performing a grid search
model = LogisticRegression()
param_grid = {
    'penalty': ['l1', 'l2', 'elasticnet'], #
    'solver' : ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag',
    ↪ 'saga'],
    'multi_class': ['auto', 'ovr', 'multinomial'],
    'max_iter' : [10,100,1000,20,200,2000]
}
grid_search = GridSearchCV(model, param_grid,cv=5, n_jobs=-1) #Grid search
    ↪ helps to get the best parameters suitable for the model
grid_search
```

```
[43]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
               param_grid={'max_iter': [10, 100, 1000, 20, 200, 2000],
                           'multi_class': ['auto', 'ovr', 'multinomial'],
                           'penalty': ['l1', 'l2', 'elasticnet'],
                           'solver': ['lbfgs', 'liblinear', 'newton-cg',
                                       'newton-cholesky', 'sag', 'saga']})
```

```
[44]: grid_search.fit(n_train, p_train)
```

```
[44]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
               param_grid={'max_iter': [10, 100, 1000, 20, 200, 2000],
                           'multi_class': ['auto', 'ovr', 'multinomial'],
                           'penalty': ['l1', 'l2', 'elasticnet'],
                           'solver': ['lbfgs', 'liblinear', 'newton-cg',
                                       'newton-cholesky', 'sag', 'saga']})
```

```
[48]: best_params=grid_search.best_params_
best_params
```

```
[48]: {'max_iter': 10, 'multi_class': 'auto', 'penalty': 'l1', 'solver': 'liblinear'}
```

```
[50]: #creating a model
best_model = LogisticRegression(**best_params)
best_model.fit(n_train,p_train)
best_model
```

```
[50]: LogisticRegression(max_iter=10, penalty='l1', solver='liblinear')
```

```
[51]: #making predictions  
p_pred = best_model.predict(n_test)  
p_pred
```

```
[51]: array([0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0],  
        dtype=int64)
```

```
[52]: #predicting the mean_absolute_error  
from sklearn.metrics import recall_score, precision_score, f1_score,  
    ↪ accuracy_score  
recall_score = recall_score(p_test,p_pred)  
precision_score = precision_score(p_test,p_pred)  
f1_score = f1_score(p_test,p_pred)  
Accuracy = accuracy_score(p_test,p_pred)  
  
print('recall_score:',recall_score )  
print('precision_score:',precision_score )  
print('f1_score:', f1_score)  
print('accuracy_score:',Accuracy)
```

```
recall_score: 1.0  
precision_score: 1.0  
f1_score: 1.0  
accuracy_score: 1.0
```

```
[ ]:
```